



# Optimized CUDA Kernels for CSR sparse matrix multiplication

CIS6010, Fall 2025

---

Anh Tran, Aditya Agarwal, Rajath Ramana  
December 08, 2025

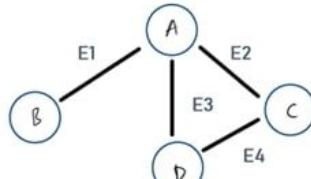
# Contents:

---

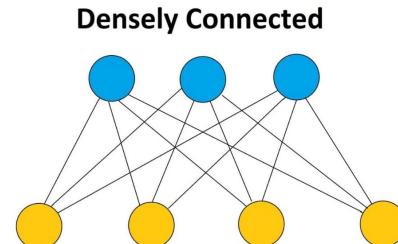
1. Problem & Motivation
2. Background
3. Methods
4. Experimental
5. Implication & Conclusion

# Motivation

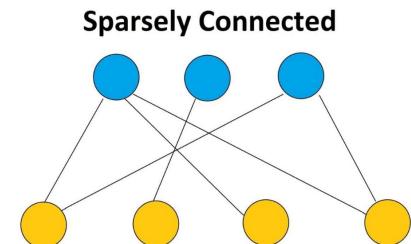
- CSR SPMM is important in many ML workloads [8] (Graph Neural Networks, Sparse Linear Layer).
- Existing Libraries (Eg. cuSPARSE [2, 3], torch.sparse) don't handle it very well.
- Need custom SPMM kernels optimised for irregular patterns.



	E1	E2	E3	E4
A	1	1	1	0
B	1	0	0	0
C	0	1	0	1
D	0	0	1	1



Densely Connected



Sparsely Connected

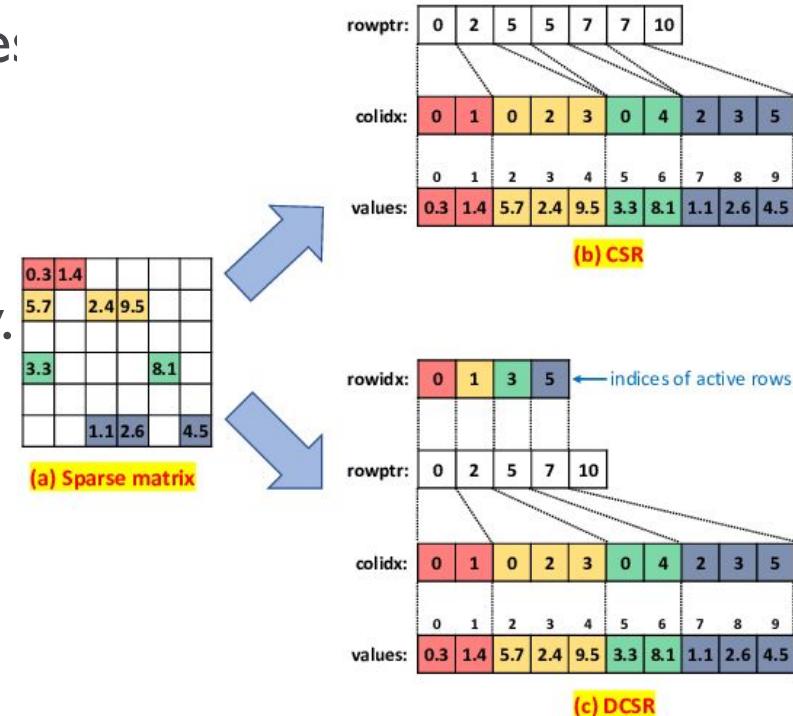
# Background

---

- CSR Format - Stores only non-zeroes using `row_ptr`, `col_idx` and `values`, this is good for unstructured sparsity.
- cuSPARSE - NVIDIA's sparse lib, slow for irregular sparse patterns.
- `torch.sparse` - Pytorch supports CSR format, relies on cuSPARSE, lacks custom kernel and speed.
- Research shows that irregular sparsities are common.

# CSR Illustration

- CSR [I] stores non-zero values, column position in col\_idx
- row\_ptr marks where each non-zero value starts in a row.
- DCSR - also has row\_idx which stores number of all non-zero rows.
- Use Case - Matrices where a lot of the rows are empty.



# Contributions of this work

---

- Highly optimized CUDA kernel, outperform cuSPARSE.
- Implemented naive, warp-per-row, shared-memory, and FP4 vectorized variants for irregular sparsity.
- Python API, outperform torch.sparse.
- Empirical Exp, analysis on performance.
- Unified benchmarks to test and compare various implementations.

# Methods

---

## Naive Approach

- Each thread is responsible for computing one output element  $C[\text{row},\text{col}]$  by taking a sparse row of A (CSR) and a full column of B.
- The thread iterates over the nonzeros in its assigned row of A and accumulates products with the corresponding elements from B.

## WarpPerRow Approach

- One warp handles one entire sparse row of A, improving locality and reducing redundant loads of that row's nonzero elements.
- Each lane in the warp computes different output columns (lane, lane+32, lane+64, ...), enabling parallel computation of the full row of C.

# Methods

Analyze workload of threads via **main loop**

$A[M, N] \times B[N, K] \Rightarrow C[M, K]$

## Naive Approach

```
sum = 0
for each nonzero entry j in row:
    A_col = column_index[j]
    A_val = value[j]
    sum += A_val * B[A_col, col]
```

## WarpPerRow Approach

```
for each assigned output column col = lane, lane+32,
lane+64, ... < K:
    sum = 0

    for each nonzero j in this sparse row:
        a = value[j]
        c = column_index[j]

        sum += a * B[c, col]

    old_value = (beta != 0) ? C[row, col] : 0
    C[row, col] = alpha * sum + beta * old_value
```

- Workload for 1 thread is typically smaller
- But loop length is arbitrary, hard to optimized

- Workload for 1 thread is heavier (nested loops), but the (outer) loop length are predictable

=> facilitate various optimization techniques:  
shared memory, packed load/store

# Methods

$$A[M, N] \times B[N, K] \Rightarrow C[M, K]$$

## Optimization strategies:

- Threads can cooperatively load A into sharedmem  
=> reduce global mem traffic
- Packed load/store: CUDA cores allow load/store a pack of **4x float32**  
=> can be applied for B and C  
=> vectorize memory accesses  
=> reduce complexity of outer loop

## WarpPerRow Approach

```
for each assigned output column col = lane, lane+32,  
lane+64, ... < K:  
    sum = 0  
  
    for each nonzero j in this sparse row:  
        a = value[j]  
        c = column_index[j]  
  
        sum += a * B[c, col]  
  
    old_value = (beta != 0) ? C[row, col] : 0  
    C[row, col] = alpha * sum + beta * old_value
```

# Experiment & Analysis

---

## Experiment Setup:

Matrix sizes tested:

- small = (512, 1024, 64)
- medium = (1024, 2048, 128)
- large = (4096, 4096, 128)
- XL = (8192, 4096, 256)

Densities: 0.01 ~ 0.30 (sprasity 0.99 ~ 0.70)

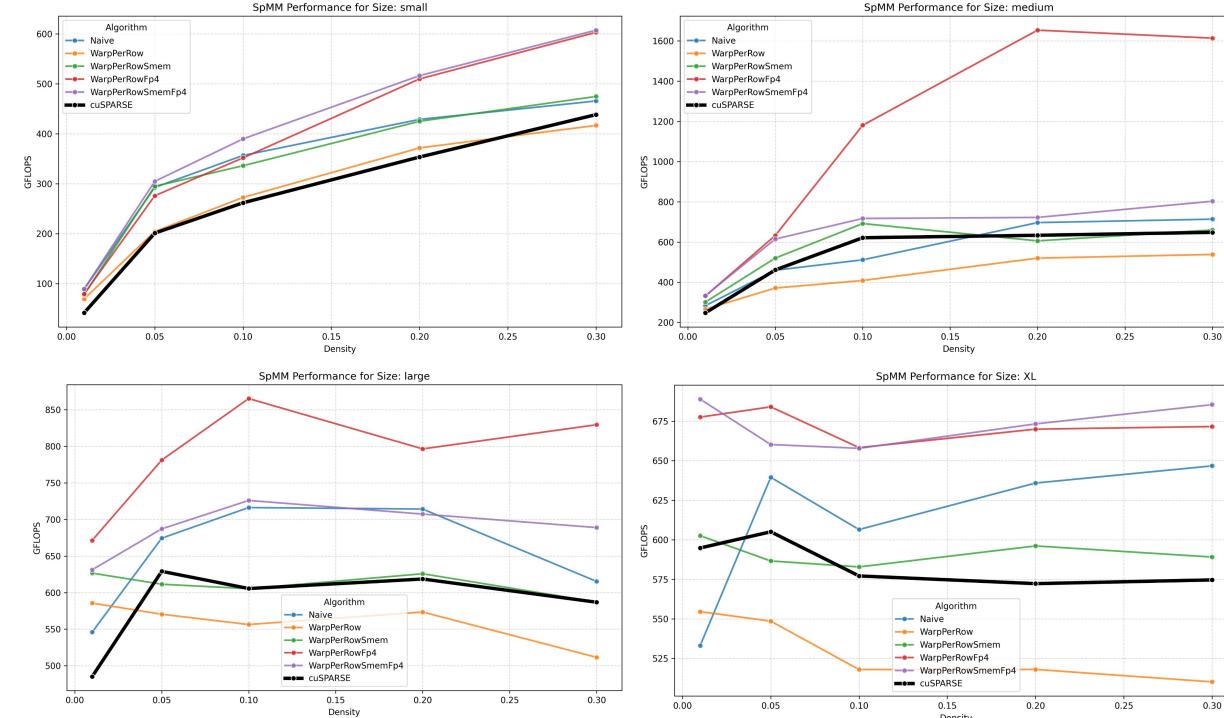
These are **typical sizes and sparsities** in Linear Layers or GNNs [5, 6, 7].

## Kernels - 5 variants

- naive
- warp per row (wpr)
- wrp + sharedmem
- wrp + fp4 (vectorized load/store 4x float32)
- wrp + sharedmem + fp4

**Hardware:** Nvidia GeForce RTX4070

# Performance vs cuSPARSE



- **WarpPerRowFp4** is consistently the top performer, with **WarpPerRowSmemFp4** is comparable on very small or very large matrices multiplication.
- Performance improvement widen with matrix size -> showing the kernels scale well.
- Our kernels dominate even at high densities -> the kernels can handle a wide sparsity range efficiently.

Note: baseline denoted by **bold, dark lines**.

# Performance vs torch.sparse

Algorithm	mean GFLOPS	mean speedup
WarpPerRowFp4	727	1.483
WarpPerRowSmemFp4	595	1.294
Naive	531	1.161
WarpPerRowSmem	520	1.134
<u>cuSPARSE</u>	<u>488</u>	<u>1.000</u>
WarpPerRow	444	0.964

*Our kernels vs cuSPARSE*

Algorithm	mean GFLOPS	mean speedup
WarpPerRowSmemFp4	479	2.297
WarpPerRowFp4	466	2.267
<u><i>torch sparse addmm</i></u>	<u>347</u>	<u>1.547</u>
WarpPerRowSmem	324	1.702
Naive	310	1.633
WarpPerRow	275	1.553
<u><i>torch sparse mm</i></u>	<u>240</u>	<u>1.000</u>

*Our Python APIs vs torch.sparse*

Our kernels outperform existing libraries:

- **1.48× faster** than cuSPARSE
- **1.38× faster** than torch.sparse.addmm
- **2.30× faster** than torch.sparse.mm

# Breakdown of Performance Factors

Metric	naive	wpr	wpr_smem	wpr_fp4	wpr_smem_fp4
GFLOPS	273.39	255.77	278.31	320.44	303.36
Memory Throughput (% peak)	94.72	86.41	76.94	73.43	73.13
Memory Bandwidth (GB/s)	12.55	11.57	12.10	13.10	11.93
L1/TEX Hit Rate (%)	73.78	33.04	2.95	18.68	3.77
L2 Hit Rate (%)	97.12	98.93	98.82	98.76	98.76
Mem Pipes Busy (%)	94.72	86.41	76.94	24.63	19.34
Coalescing Efficiency ( / 32)	22.7	22.7	31.9	28.9	31.9
Warp Cycles / Instruction	29.02	29.98	16.12	68.66	52.07
Achieved Occupancy (%)	63.54	91.68	46.84	77.57	46.94
Active Warps per SM	30.50	44.0	22.48	37.23	22.53
IPC (inst / cycle)	1.05	1.47	1.39	0.54	0.43
Total Instructions Executed	178,205,120	267,000,000	243,099,537	88,314,836	76,548,614

## Better coalescing

- **Shared memory** allows us to restructure the matrix layout and access pattern, to fit the warps better
- Coalescing efficiency improved significantly with shared mem

Profiling data. Size [8192 x 4096 x 256]. Density 0.1

# Breakdown of Performance Factors

Metric	naive	wpr	wpr_smem	wpr_fp4	wpr_smem_fp4
GFLOPS	273.39	255.77	278.31	320.44	303.36
Memory Throughput (% peak)	94.72	86.41	76.94	73.43	73.13
Memory Bandwidth (GB/s)	12.55	11.57	12.10	13.10	11.93
L1/TEX Hit Rate (%)	73.78	33.04	2.95	18.68	3.77
L2 Hit Rate (%)	97.12	98.93	98.82	98.76	98.76
Mem Pipes Busy (%)	94.72	86.41	76.94	24.63	19.34
Coalescing Efficiency (/ 32)	22.7	22.7	31.9	28.9	31.9
Warp Cycles / Instruction	29.02	29.98	16.12	68.66	52.07
Achieved Occupancy (%)	63.54	91.68	46.84	77.57	46.94
Active Warps per SM	30.50	44.0	22.48	37.23	22.53
IPC (inst / cycle)	1.05	1.47	1.39	0.54	0.43
Total Instructions Executed	178,205,120	267,000,000	243,099,537	88,314,836	76,548,614

Profiling data. Size [8192 x 4096 x 256]. Density 0.1

Reduce pressure on memory

- Naive/WPR are ~94% (heavily memory-bound). Shared-memory cuts this to ~77%, and FP4 variants drop to ~20–25%, a 4-5x reduction in memory load.
- FP4 packing + shared-memory reuse move the kernel from memory-bound to compute-bound.

# Breakdown of Performance Factors

Metric	naive	wpr	wpr_smem	wpr_fp4	wpr_smem_fp4
GFLOPS	273.39	255.77	278.31	320.44	303.36
Memory Throughput (% peak)	94.72	86.41	76.94	73.43	73.13
Memory Bandwidth (GB/s)	12.55	11.57	12.10	13.10	11.93
L1/TEX Hit Rate (%)	73.78	33.04	2.95	18.68	3.77
L2 Hit Rate (%)	97.12	98.93	98.82	98.76	98.76
Mem Pipes Busy (%)	94.72	86.41	76.94	24.63	19.34
Coalescing Efficiency ( / 32)	22.7	22.7	31.9	28.9	31.9
Warp Cycles / Instruction	29.02	29.98	16.12	68.66	52.07
Achieved Occupancy (%)	63.54	91.68	46.84	77.57	46.94
Active Warps per SM	30.50	44.0	22.48	37.23	22.53
IPC (inst / cycle)	1.05	1.47	1.39	0.54	0.43
Total Instructions Executed	178,205,120	267,000,000	243,099,537	88,314,836	76,548,614

## Fewer instructions

- SMEM slightly reduce # of instructions to 243M and 76M (fewer global-access insns)
- FP4 packing lowers total instructions from 178M (naive) to 88M and 76M (FP4 variants)  
**(more aggressive; basically, we shorten the loop)**

Profiling data. Size [8192 x 4096 x 256]. Density 0.1

# Breakdown of Performance Factors

Metric	naive	wpr	wpr_smem	wpr_fp4	wpr_smem_fp4
GFLOPS	273.39	255.77	278.31	320.44	303.36
Memory Throughput (% peak)	94.72	86.41	76.94	73.43	73.13
Memory Bandwidth (GB/s)	12.55	11.57	12.10	13.10	11.93
L1/TEX Hit Rate (%)	73.78	33.04	2.95	18.68	3.77
L2 Hit Rate (%)	97.12	98.93	98.82	98.76	98.76
Mem Pipes Busy (%)	94.72	86.41	76.94	24.63	19.34
Coalescing Efficiency ( / 32)	22.7	22.7	31.9	28.9	31.9
Warp Cycles / Instruction	29.02	29.98	16.12	68.66	52.07
Achieved Occupancy (%)	63.54	91.68	46.84	77.57	46.94
Active Warps per SM	30.50	44.0	22.48	37.23	22.53
IPC (inst / cycle)	1.05	1.47	1.39	0.54	0.43
Total Instructions Executed	178,205,120	267,000,000	243,099,537	88,314,836	76,548,614

Profiling data. Size [8192 x 4096 x 256]. Density 0.1

One things that we are not (yet) able to fully explain is **cache hit rates**.

Share mem + vectorized load/store (Fp4) has low hit rates but, somehow, higher throughput.

## Our hypothesis:

- many threads in naive kernel loads the same cols/rows  
=> more frequent accesses  
=> explain higher rates  
=> naive kernels might not benefit a lot from this
- significantly lower Memory pipe busy and total instructions of smem and fp4 compensate the poor cache rates.

# Conclusions

---

- Developed a clean, modular CSR-SpMM library with multiple CUDA kernel variants.
- Warp-optimized kernels **outperform cuSPARSE and PyTorch sparse** on many sparsity regimes.
- Provided a simple PyTorch API and benchmarking tools for easy evaluation and extension.
- Demonstrated clear performance trade-offs and a solid foundation for future kernel research.

# Shortcomings and Future Scope

---

- Limitations:
  - Handles general CSR matrices but does not yet optimize for extreme irregular sparsity patterns.
  - Dynamic warp scheduling was explored but not fully integrated or validated.
- Future scope:
  - Add advanced performance features (streaming, better tiling, load balancing) and broaden Python/CUDA interfaces for more sparsity formats and end-to-end use.

# Work Distribution

---

Anh: Developed 5 kernels, debugged and integrated Python APIs, designed benchmarks/experiments, analyzed results.

Aditya & Rajath: Implemented the PyTorch integration stack/interface and contributed to developing the naïve shared-memory kernel.

All team members: Collaborated on benchmarking, verification, and end-to-end system integration, comparing the custom CSR-SpMM kernels with PyTorch's sparse operations.

# Reference

---

- [1] Matt Eding, “*Sparse Matrices*” — a comprehensive overview of sparse-matrix data structures and formats (COO, CSR/CSC, DOK, LIL, BSR, DIA), with pros/cons for each. [matteding.github.io](https://matteding.github.io)
- [2] NVIDIA Developer Blog: “*Accelerating Inference with Sparsity Using the NVIDIA Ampere Architecture and NVIDIA TensorRT*” — describes how TensorRT 8.0 + Ampere-GPU Sparse Tensor Cores leverage 2:4 structured sparsity for ~50% zeros and deliver large performance and efficiency gains for neural-network inference. [NVIDIA Developer](#)
- [3] NVPL SPARSE documentation: “*Sparse Matrix Formats*” — outlines storage formats (COO, CSR, CSC, etc.) supported by NVPL, with detailed definitions of how sparse matrices are represented in memory. [NVIDIA Docs](#)
- [4] cuSPARSE Generic API documentation — reference guide for generic sparse-matrix operations under NVIDIA’s CUDA ecosystem. [NVIDIA Docs](#)

# Reference

---

- [5] Jonathan Frankle and Michael Carbin, The Lottery Ticket Hypothesis: Finding Sparse, Trainable Neural Networks, published at ICLR 2019.
- [6] Victor Sanh, Thomas Wolf, and Alexander M. Rush, Movement Pruning: Adaptive Sparsity by Fine-Tuning, published at NeurIPS 2020.
- [7] Tianlong Chen, Yongduo Sui, Xuxi Chen, Aston Zhang, and Zhangyang Wang, A Unified Lottery Ticket Hypothesis for Graph Neural Networks, published at ICML 2021.
- [8] Tasou, I., Mpakos, P., Vichos, A., et al. Sparse Computations in Deep Learning Inference. arXiv:2512.02550, 2025.